

Assignment 0

CSE 130: Principles of Computer System Design, Fall 2019

Due: April 15 at 9:00PM

Goals

There are three goals for Assignment 0. The first is to get your programming environment set up properly, which will require that you install an Ubuntu 18.04 virtual machine on your personal computer. The second is to learn the lab format for CSE 130 by writing a simple program that acts similar to the UNIX `cat` command with a simple modification. As with all programming assignments, this lab will require that you submit a design document, a README file, and a lab writeup along with your code in your `git` repository. The third goal is to ensure that you're ready for this class by giving you a straightforward assignment that will act as a "self-test". If you have a lot of difficulty with this assignment, you may not be ready for CSE 130.

Setting up Ubuntu

Instructions on how to set up an Ubuntu 18.04 VM are available on Piazza.

Programming assignment: `dog`

Design document

Before writing code for this assignment, as with every other assignment, you must write up a design document. Your design document must be called `DESIGN.pdf`, and must be in PDF (you can easily convert other document formats, including plain text and Markdown, to PDF).

Your design document should describe the design of your code in sufficient detail that a knowledgeable programmer could duplicate your work. This includes descriptions of the data structures you use, all non-trivial algorithms and formulas, and a description of each function including its purpose, inputs, outputs, and assumptions it makes about the inputs or outputs. A sample design document is available from the resources page on Canvas.

Yes, the design for `dog` will be short, but we want you to get experience writing one up on a simple program before tackling more difficult programs. **Your design document is a significant fraction of the grade for each assignment**, so get in the habit of writing a good design document *before* you start writing code. It'll make writing code a lot easier. Also, if you want help with your code, the first thing we're going to ask for is your design document. We're happy to help you with the design, but we can't debug code without a design any more than you can.

Program functionality

The only code you have to write for this assignment is to implement the basic `cat` program, *without support for any flags and with file arguments in reverse order*. That means your code needs to copy data from each of the files specified on the command line to standard output, *in the reverse order that they appear in the command line*. For example, `dog file1 file2 file3` will copy all of the data from `file3`, `file2`, and `file1` to standard output, in that order, same as calling `cat file3 file2 file1`. If `-` (dash) is given as a filename, `dog` uses standard input for that file (not a file named `-`). Note that `-` may only be given as one file name, just as `cat` allows.

Your program is called `dog` (rather than `cat`) to ensure that when you run it, you don't accidentally run the installed version of `cat`.

If no files are specified on the command line, `dog` should just copy standard input to standard output until it runs out of input, just like the installed version of `cat` does. Note that the data might be binary; your code must work in that case, that is the output must be the same as if you were to run the command `cat` with the corresponding arguments. Process files one at a time; if `dog` runs into an error with a file, the program should print an error message to standard error (not standard output!) and skip the file, handling the remainder of the files. Your error messages should be identical (except for the program name) to those printed by `cat`; the `strerror(3)` library function will be useful for this.

Your program may not use any of the C library `FILE` * functions such as `fread()` and `printf()` for user data. C++ library functions such as `iostream` and `fstream` are also not allowed. You may use `fprintf()`, `perror()`, and `warn()` for error messages, but you may not use `fopen()`. `sprintf()` is fine—it doesn't take a `FILE *` as an argument. Your code must use fixed-size buffers, and may allocate no more than 32 KiB of memory for them (either via `malloc()` or as a direct variable declaration).

Your code must be in C. The same will be true for all future assignments: they must be written in C. Please familiarize yourself with the Coding Guide for this class, available on Canvas. Following good coding practices will not be graded directly, but can help reduce the amount of bugs and lets the course staff help you more efficiently.

Testing your code

You should test your code on your own system. Make up commands, and try them on both `cat` and `dog`. Use `diff(1)` to see if the commands work the same.

You might also consider cloning a new copy of your repository (from `GITLAB@UCSC`) to a clean directory to see if it builds properly, and runs as you expect. That's an easy way to tell if your repository has all of the right files in it. You can then delete the newly-cloned copy of the directory on your local machine once you're done with it.

We plan to provide a service on `GITLAB@UCSC` that will allow you to run a *subset* of the tests that we'll run on your code for each assignment. You will be able to run this test from

the GITLAB@UCSC server at most twice per day (days start at midnight), and (of course) you can only run it on commits that have been pushed to GITLAB@UCSC. Running these tests is completely optional, and we'll go over how to do it in section during the first week of class.

The GITLAB@UCSC test will cover at least half of the functionality points for this assignment, **but there will be additional test cases not covered by this service**, so you should still do your own testing.

README and Writeup

Your repository must include a README file in Markdown format (see <https://www.markdownguide.org> for details) that *must* be named README.md. Markdown is a simple markup language that provides annotations for (plain ASCII) text to be shown in bold, italics, section headers, etc. A plain ASCII text file is a valid Markdown document—you aren't *required* to use any Markdown annotations.

The README.md file should be short, and contain any instructions necessary for running your code. You should also list limitations or issues in README.md, telling a user if there are any known issues with your code.

Your WRITEUP.pdf is where you'll describe the testing you did on your program and answer any short questions the assignment might ask. The testing can be unit testing (testing of individual functions or smaller pieces of the program) or whole-system testing, which involves running your code in particular scenarios. WRITEUP.pdf *must* have that name, and must be PDF—there are many ways to convert plain text, Word documents, and Markdown into PDF.

For Assignment 0, your writeup must answer the following question:

- How does the code for handling a file differ from that for handling standard input? What concept is this an example of?

Submitting your assignment

All of your files for Assignment 0 must be in the asgn0 directory in your git repository. Each time you push your repository to GITLAB@UCSC, the server will run a program to check the following:

- There are no “bad” files in the asgn0 directory (*i.e.*, object files)
- Your assignment builds properly in asgn0 using make to produce a dog binary
- All required files (DESIGN.pdf, README.md, WRITEUP.pdf) are present in asgn0

If the repository meets these minimum requirements for Assignment 0, there will be a green check next to your commit ID in the GITLAB@UCSC Web GUI. If it doesn't, there will be a red X. **It's OK to commit and push a repository that doesn't meet minimum requirements**

for grading—we will not take off points for doing so. However, we will only *grade* a commit that meets these minimum requirements, and has a green check mark associated with it.

Note that the *minimum* requirements say nothing about correct functionality—the green check only means that the system successfully ran `make` and that all of the required files were present, with the correct names.

You must submit the commit ID you want us to grade via Google Form, linked to the assignment page on Canvas. This must be done before the assignment deadline.

Hints

- This is a straightforward assignment, designed to get you familiar with the tools you'll need for the remaining assignments. Most of your time will be spent setting up the VM, which isn't difficult, but can take a few hours. Once you have the VM set up (or even while you're setting it up), do your design document *before* writing any code.
- Become familiar with the Coding Guide for this class, available on Canvas. Do this *before* writing your code, so that you don't need to go back and rewrite it. In addition, become familiar with the workflow you'll be following when writing code for this class, also available from Canvas.
- Your commit must contain the following files:
 - `README.md`
 - `DESIGN.pdf`
 - `Makefile`
 - `dog.c`
 - `WRITEUP.pdf`

It may *not* contain any `.o` files. You may, if you wish, include the “source” files for your `DESIGN.pdf` and/or `WRITEUP.pdf` in your repo, but you don't have to.

- You can create a zero-byte file using the `touch` command. `touch README.md` creates a zero-byte `README.md` file if none already exists. If it does exist, it does nothing besides accessing the file, which is harmless and non-destructive.
- If you need help, use online documentation such as `man` pages and documentation on `Makefiles`. If you still need help, ask the course staff. You should be familiar with the rules on academic integrity *before* you start the assignment.

Grading

As with all of the assignments in this class, we will be grading you on *all* of the material you turn in, with the approximate distribution of points as follows: design document (20%); functionality (70%); writeup (10%).

If you submit a commit ID without a green checkmark next to it or modify `.gitlab-ci.yml` in any way, your maximum grade for this assignment is 5%. Make sure the commit ID you submit on the Google form is one that has a green checkmark.